

BISCAD technical handbook

The CAD kernel for agents. build123d / OpenCascade behind a REST API and an MCP server, with a web viewer, a studio, and outputs built for humans to oversee what agents make.

This handbook takes you from zero to deep understanding of the whole system. It explains every component, the theory behind it, the design decisions and their trade-offs, and it links to the code. Read it top to bottom once. After that, use the table of contents to jump.

Prose follows `.construction/how-to-write-prose.md`: short active sentences, concrete nouns, units on numbers. Code citations use file paths, not line numbers, because the server modules are being renamed (behaviour unchanged) while this is written.



Contents

1. [Why BISCAD exists](#)
2. [System architecture](#)
3. [Geometry kernel primer](#)
4. [Stable references and the topological naming problem](#)
5. [Tessellation and the scene format](#)
6. [Headless rendering](#)
7. [Sandbox security](#)
8. [Build-step capture](#)
9. [Analysis algorithms](#)
10. [Exports, imports and drawings](#)
11. [Text-to-CAD agent loop and MCP tool design](#)
12. [Web viewer and studio internals](#)
13. [Data model, accounts, quotas, versions](#)
14. [Performance: measured numbers](#)
15. [Deployment and scaling](#)
16. [Testing strategy](#)
17. [Roadmap](#)
18. [Glossary](#)
19. [Reading list](#)

1. Why BISCAD exists

The Onshape API wall

Onshape is a cloud CAD system with a public REST API. The API is good. The limit on how much you may call it is the problem.

Onshape caps API calls per user per year. The documented limits are:

Plan	Annual API call limit
Free / Standard / EDU student	2,500 per user
Professional	5,000 per user
Enterprise	10,000 per full user

Source: the Onshape API documentation on limits (onshape-public.github.io/docs/auth/limits). When a user passes the annual limit, Onshape rejects further calls with HTTP 402 (Payment Required). The pricing page (onshape.com/pricing) lists Standard at \$1,500 per user per year and Professional at \$2,500 per user per year. So a Professional seat costs \$2,500 per year and grants 5,000 API calls in that year.

Count what a call is. Reading a document is a call. Starting a feature is a call. Reading back the result is a call. A single "make this change and tell me what it did" round trip is several calls.

Agents do not iterate like people

A person edits a CAD model a few times a minute. An agent does not. An agent writes code, builds it, looks at the result, finds the hole is in the wrong place, moves the hole, builds again. It does this in a tight loop, as fast as the kernel answers.

A text-to-CAD agent run in BISCAD is 2 to 4 build rounds. A session that explores a design space is hundreds of builds per hour. At that rate a Professional Onshape seat's whole annual allowance of 5,000 calls is gone in an afternoon. The pricing model assumes a human at a mouse. It breaks the moment the caller is a program.

BISCAD is priced per build instead of per seat. A build is a few hundred milliseconds of one CPU core (Section 14). On autoscaling infrastructure that is a fraction of a cent (Section 15). An agent can iterate hundreds of times an hour for the cost of a coffee, not a software licence.

Code as CAD, with build123d

BISCAD does not try to clone Onshape's feature tree UI. It takes a different base: **the model is a program**.

[build123d](#) is a Python CAD library on top of OpenCascade (OCCT), the same B-rep kernel that FreeCAD and many commercial tools use. A part is Python code:

```
from build123d import *
with BuildPart() as bracket:
    Box(60, 40, 6)
    fillet(bracket.edges().filter_by(Axis.Z), 3)
    with GridLocations(40, 24, 2, 2):
        Hole(3.25)
result = bracket.part
```

Code as the source of truth has properties agents love:

- It is text. A language model writes it, reads it, diffs it, and reasons about it.
- It is parametric by construction. A variable is a parameter. No special UI needed.
- It is deterministic. The same program and the same inputs give the same solid.
- It is versionable with ordinary tools.

The trade-off is that the user must think in code, not in direct manipulation. For a human that is a cost. For an agent it is the native tongue.

Humans still need to understand the output

Andrej Karpathy's point about working with large language models is that the hard part is not getting the model to produce output. It is the human staying in the loop: checking, trusting, and correcting what the model made. Output you cannot inspect is

output you cannot trust.

A wall of geometry numbers is not inspectable. So BISCAD spends real effort on making an agent's work legible to a person:

- **Renders the agent can see** (Section 6): any view, with face-id labels, so a multimodal model checks its own geometry and a person checks it too.
- **Stable references** (Section 4): every face and edge has a name like `p0/f3` that is the same in the viewer, the REST API and the MCP tools. A person and an agent point at the same face with the same word.
- **Build-step replay** (Section 8): one plain-English line per operation, with a thumbnail, so you read how the part was built the way you read a recipe.
- **Design reports and explainer GIFs** (Section 11, `server/explainer.py`): one page that shows the part, the steps, the mass, and the manufacturability check.

The thesis of BISCAD: an AI makes the geometry, the human oversees it, and the tooling exists to make that oversight fast.

2. System architecture

Components

```
graph TD
  subgraph Clients
    A1[Agent via MCP]
    A2[Browser studio / viewer]
    A3[CLI / Python SDK<br/>client/voncad.py]
  end
  subgraph Server["FastAPI process (server/app.py)"]
    REST[REST routes<br/>/v1/*]
    MCP[MCP server<br/>/mcp · mcp_server.py]
    CORE[core.py<br/>identity · quotas · rate limit]
    STORE[store.py<br/>SQLite + files]
  end
  subgraph Kernel["Forked worker per job (kernel.py)"]
    GUARD[AST guard + sanitized modules]
    B123[build123d / OpenCascade]
    TESS[tessellate.py]
    STEPS[steps.py]
  end
  A1 --> MCP
  A2 --> REST
  A3 --> REST
  MCP --> CORE
  REST --> CORE
  CORE --> STORE
  CORE -->|fork| Kernel
  Kernel -->|scene.json · mesh.npz · *.brep| STORE
  REST -. render/analysis .-> Kernel
```

The whole backend is one FastAPI process. It serves three faces of the same core logic:

- **REST** (`server/app.py`): the HTTP API. Contract in `README.md` / `PROGRESS.md`.
- **MCP** (`server/mcp_server.py`): the same operations as tools an agent calls, mounted at `/mcp`. Renders come back as real images.
- **Static web** (`web/`): the landing page, studio, embeddable viewer, and docs.

Shared business logic lives in `server/core.py` so REST and MCP never drift apart. Both call `core.build`, `core.measure`, and so on. Identity, rate limiting and quotas are decided once, in `core.py`, for every caller.

The kernel is isolated. Every build and every analysis runs in a forked child process with resource limits (Section 7). The parent process never runs untrusted code in its own interpreter.

The lifecycle of POST `/v1/build`

```
sequenceDiagram
    participant C as Client
    participant App as app.py (REST)
    participant Core as core.py
    participant Store as store.py
    participant Child as forked worker (kernel.py)
    participant FS as /data/versions/<vid>

    C->>App: POST /v1/build {script, params, quality}
    App->>Core: identify(auth, ip) · rate_limit · check_build_quota
    Note over Core: 401 / 402 / 429 here if the caller is over a limit
    App->>Core: core.build(...)
    Core->>Core: check_script(src) (fast AST reject, no fork)
    Core->>Store: new_id("v") → vid
    Core->>Child: fork() + run kernel._build_job in child
    Note over Child: RLIMIT_AS/CPU/FSIZE set<br/>FSIZE=0 while user code runs
    Child->>Child: exec script → build123d solids
    Child->>Child: tessellate → scene, mesh, topology
    Child->>Child: capture build steps (draft mesh each)
    Child->>FS: scene.json · topology.json · mesh.npz · part*.brep · steps/
    Child->>Core: summary + logs + param_schema (via pipe)
    Core->>Store: save_version(...) in SQLite
    Core->>App: public_version(v) with links
    App-->>C: {ok, summary, scene?, steps, links}
```

Steps worth noting:

1. **Cheap reject first.** `core.build` calls the AST guard before it forks. A script that imports `os` fails in microseconds, with no process spawned.
2. **One capability id.** The version id (`v_...`) is minted before the build. It names both the SQLite row and the on-disk directory `data/versions/<vid>/.`
3. **Fork, not spawn.** `build123d` and `OpenCascade` are already imported in the parent. `fork()` copies that warm memory, so the child starts ready. Spawning a fresh interpreter would re-import OCCT and cost over 2 seconds (measured; Section 14).
4. **Files are the build output.** The child writes `scene.json` (what the viewer draws), `topology.json` (what an agent reads), `mesh.npz` (what the renderer uses), and `part*.brep` (the exact B-rep for later exports and analysis). The parent never holds the geometry; it reads these files on demand.
5. **The result is a capability link.** The response carries URLs for the scene, render, report, exports and viewer. Anyone with the version id can view it (Section 13).

Render and analysis requests follow the same pattern: the parent forks a child that loads the saved `.brep`, does the work, and returns JSON or an image.

3. Geometry kernel primer

To understand everything above the kernel, you need to understand the kernel's model of shape. This section is the CAD theory. If you know OCCT already, skim it.

Three ways to represent a solid

- **Mesh (boundary triangles).** A bag of triangles approximating the surface. Fast to draw, cheap to store, lossy. An STL file is a mesh. A mesh does not know it has a "hole" or a "cylinder"; it only has triangles. You cannot fillet a mesh edge exactly, because the mesh has no exact edge, only facets.
- **CSG (constructive solid geometry).** A tree of primitives combined with booleans: "box minus cylinder union sphere". OpenSCAD works this way. CSG is exact and simple to author, but it has no persistent notion of a face. "The top face" is not a thing the tree stores; it is recomputed every time.
- **B-rep (boundary representation).** The surface is stored as exact trimmed surfaces (planes, cylinders, NURBS) bounded by exact curves, stitched into a watertight shell. A B-rep knows its faces, edges and vertices as first-class objects with identity and exact geometry. This is what commercial CAD uses, and it is what OpenCascade is.

BISCAD is B-rep. Faces and edges are real objects you can name, measure, fillet and select. The mesh (Section 5) is a derived view for drawing; the `.brep` file is the truth.

OpenCascade topology: TopoDS

OpenCascade (OCCT) splits a shape into **geometry** (the math of a surface or curve) and **topology** (how pieces connect). Topology is the `TopoDS_Shape` hierarchy:

```
Compound → CompSolid → Solid → Shell → Face → Wire → Edge → Vertex
```

- A **Face** is a trimmed patch of one surface (a plane, a cylinder, a NURBS surface).
- A **Wire** is a loop of edges bounding a face. A face has one outer wire and zero or more inner wires (holes in the face).
- An **Edge** is a trimmed curve (a line, an arc, a spline) shared by two faces.
- A **Vertex** is a point where edges meet.

Two faces share one edge. The edge is stored once and referenced twice. This sharing is what makes a B-rep watertight and is why topology and geometry are separate: the same curve geometry can be referenced with either orientation by the two faces that meet along it.

Orientation matters. Each face carries an orientation flag (`FORWARD` or `REVERSED`) that tells you which side is "outside" the solid. The tessellator reads this flag to wind triangles consistently and to flip surface normals so they point outward (`tessellate.py`, `_mesh_face`, checks `TopAbs_REVERSED`). Get it wrong and the part renders inside-out.

build123d: builder mode vs algebra mode

build123d wraps OCCT in Pythonic form, and offers two authoring styles.

Builder mode uses a `with` context that collects operations into a running part:

```
with BuildPart() as p:
    Box(40, 30, 10)
    Cylinder(5, 20, mode=Mode.SUBTRACT)
    fillet(p.edges().filter_by(Axis.Z), 2)
result = p.part
```

Operations like `extrude`, `Hole` and `fillet` add to the context implicitly. This is the style BISCAD's step capture hooks into (Section 8), and it is what the text-to-CAD agent is told to prefer, because it produces readable build steps.

Algebra mode uses plain operators on shape objects:

```
result = Box(40, 30, 10) - Cylinder(5, 20)
```

Algebra mode is terse and functional. It produces no builder context, so BISCAD records no step replay for it; the design report says so ("built in algebra mode"). Both modes produce the same B-rep.

Booleans, fillets, chamfers, and why they fail

Booleans (union, cut, intersect) are the core B-rep operation. OCCT's `BRepAlgoAPI` intersects the two shells, finds the intersection curves, trims the faces, and sews a new watertight shell. This is the single most expensive and most fragile operation in the kernel. It is the top cost in a heavy build (Section 14: `BRepAlgoAPI` was the largest single time sink in the profiled quadruped build).

Fillets round an edge by replacing it with a blend surface (a rolling-ball surface of the given radius). **Chamfers** replace it with a flat bevel. Both are local surgery on the B-rep: they remove the edge, build the new surface, and re-trim the neighbours.

They fail for real geometric reasons, not bugs:

- **Radius too large.** A fillet radius larger than the local feature leaves no room for the blend. The rolling ball self-intersects. OCCT throws, and BISCAD reports the line.
- **Convergent blends.** Fillets on edges that meet at a vertex must blend into each other. If three fillets of different radii meet, the corner patch may be unsolvable.
- **Tangent or near-tangent faces.** When the two faces meeting at an edge are nearly tangent, the blend surface is ill-conditioned.

Because they fail, BISCAD treats a failed build as a normal event. The error text and the failing script line go back to the caller (and to the agent, which then shrinks the radius and retries).

Tolerance and validity

A B-rep is not infinitely precise. Two surfaces that should meet exactly meet within a **tolerance** (typically $1e-7$ to $1e-3$ mm per vertex/edge/face). Tolerance accumulates through booleans. A shape whose tolerances have grown too large, or whose faces no longer sew into a closed shell, is **invalid**.

BISCAD checks validity with OCCT's `BRepCheck` (via `shape.is_valid()`) on every part and reports it in the summary (`kernel.py`, `summarize_parts`) and in the DFM check (`analysis.py`). An invalid solid will not export to STEP cleanly and may crash a downstream tool, so catching it early matters. Validity checking is itself expensive: it was the second-largest cost in the profiled build (Section 14).

4. Stable references and the topological naming problem

How p0/f3 is assigned


Every face, edge and vertex in BISCAD has a stable-looking name:

```
p0/f3 = part 0, face index 3
p0/e7 = part 0, edge index 7
p0/v2 = part 0, vertex index 2
```

The index is the position of that entity in `build123d's shape.faces() / shape.edges() / shape.vertices()` list. `build123d` builds those lists by walking the OCCT topology with `TopExp_Explorer`, which visits sub-shapes in a deterministic order for a given shape. So `p0/f3` is exactly `result.faces()[3]` for a single part.

This is the whole point: one naming scheme, used everywhere.

- A script selects `part.faces()[3]`.
- The REST topology endpoint returns `{"id": "p0/f3", ...}`.
- The viewer's GPU picker returns `p0/f3` when you click that face.
- An agent says "fillet `p0/f3`" and the human sees the same face highlighted.



The assignment lives in `tessellate.py` (`mesh_shape` writes `f"{pid}/f{k}"` as it enumerates faces) and is read back in `analysis.py` (`_entity` parses `p0/f3` back to `shape.faces()[3]`).

When the ids are stable

The ids are stable **for the same program and the same parameters**. Build the same script twice, get the same ids. This is enough for the common loop: an agent builds, reads topology, measures `p0/f3`, sections at `p0/f3`, exports. Nothing changes between those calls because the geometry does not change.

Ids are also stable under changes that do not alter the face **count or order**. A pure dimensional change is the good case. Measured: the angle bracket built at `width=60` and at `width=80` keeps all 18 faces in the same order and type (verified directly: 18

of 18 faces match type and normal). The holes move, but `p0/f3` is still the same face.


When the ids are NOT stable: the topological naming problem

Add a feature and the ids shift. This is the **topological naming problem** (TNP), the oldest unsolved problem in parametric B-rep CAD.


The problem: a B-rep face has no permanent identity. It is identified only by its position in a list that is rebuilt from scratch every time the model regenerates. When an earlier operation changes how many faces exist, or in what order, every later reference by index points at a different face.

Measured example. Take a plate with holes:

Program	Faces	"Top face" (normal +Z)
holes=2, fillet=0	8	p0/f2
holes=3, fillet=0	9	p0/f2 (survived: holes added at the end)
holes=2, fillet=3	12	p0/f1 — and old p0/f2 is now a fillet face



holes=2, fillet=0: top face is f2



holes=2, fillet=3: top face is f1, f2 is a fillet

Adding the fillet inserted new faces earlier in the walk order, so every later index slid. A reference to `p0/f2` that meant "the top face" now means "a fillet face". An agent that stored `p0/f2` across an edit would fillet the wrong thing.

This is not a BISCAD bug. It is intrinsic to naming faces by regeneration order. BISCAD's honest position: ids are stable within a build, and stable across pure parameter changes, and you must re-read topology after a structural edit. The MCP guide text and the docs say exactly this ("Ids are stable for the same program and parameters").

How other systems attack TNP

- **Onshape** solved it most completely, with a system often described under the name *Topological Naming/deterministic ids*. Onshape assigns each created entity a persistent id at the moment the feature creates it, derived from the operation and its inputs, and tracks those ids through later operations. A face keeps its id across edits because the id is stored, not recomputed from position. This is a major reason Onshape's parametric edits are robust, and a major reason its system is proprietary and complex.
- **FreeCAD** suffered TNP badly for years (references breaking on edit was the top user complaint). The long-running "Toponaming" fix, merged for FreeCAD 1.0, assigns and tracks string names for sub-shapes through operations, mapping new sub-shapes back to the elements that generated them. It is based on Realthunder's element-mapping work.
- **Parasolid / ACIS** (commercial kernels) carry attribute/rollback machinery that tags entities through the history so references survive.

Future work for BISCAD

The path to persistent naming in BISCAD is clear and does not need a new kernel:

1. **Record a build history.** The step recorder (Section 8) already wraps every operation. Extend it to tag the faces each operation creates.

2. **Use OCCT's** `BRepTools_History` / `ShapeHistory`. OCCT's boolean and feature operations can return a history object that maps input sub-shapes to the output sub-shapes they became (generated / modified / deleted). Threading that through gives each face a lineage: "the top face of the box, as modified by the fillet".
3. **Expose a persistent id alongside the index id.** Keep `p0/f3` for the within-build case, and add `p0/#<stable>` that survives edits. The viewer and MCP already carry arbitrary id strings, so the plumbing is in place.

This is the single highest-value piece of future work, because robust cross-edit references are what let an agent refine a design over many rounds without re-reading everything each time.

5. Tessellation and the scene format

The viewer and the renderer need triangles, not B-rep. Tessellation turns the exact surface into a mesh. It lives in `server/tessellate.py`.

BRepMesh: the incremental mesher

OCCT's `BRepMesh_IncrementalMesh` meshes a shape face by face. Two parameters control how closely the triangles follow the surface:

- **Linear deflection** (mm): the maximum distance between the mesh and the true surface. Smaller deflection, more triangles, closer fit.
- **Angular deflection** (radians): the maximum angle between adjacent facet normals. This controls how many facets go around a curve. It is why a small deflection still gives a round-looking cylinder.

BISCAD scales the linear deflection to the model's size so a big part and a small part look equally smooth. `tolerance_for(shape)` computes `deflection = diagonal / divisor`, clamped to a 0.002 mm floor, where the divisor comes from the quality preset.


Quality presets

```
QUALITY = {"draft": (600.0, 0.5), "normal": (2000.0, 0.25), "fine": (6000.0, 0.12)}
#           divisor  angle(rad)
```

The number is "diagonal / divisor". A larger divisor means a finer mesh. The angle tightens with quality.

The effect, measured on the flange (one part, 27 faces):

Quality	Triangles	scene.json
draft	834	136 KB
normal	1,926	267 KB
fine	5,128	628 KB



The studio builds at `normal`. Build-step snapshots are meshed at `draft`, because they are previews and there may be dozens of them (Section 8). You pick quality per build via the API (`quality` field) or the studio's quality selector.

Per-face triangle ranges

This is the detail that makes GPU picking work (Section 12). When BISCAD meshes a part, it concatenates every face's triangles into one buffer **in face order**, and records for each face the range it occupies:

```
{"id": "p0/f3", "type": "plane", "area": 2743.27,
 "center": [29.81, 0.0, 0.0], "normal": [0,0,-1],
 "start": 164, "count": 216}
```

`start` is the index of the first triangle, `count` is how many. Triangles of one face are contiguous, so a picked triangle number maps to its face by a range lookup. This is why clicking a triangle in the viewer gives you `p0/f3` and not just "some triangle". (Verified:


for the bracket, the face ranges sum exactly to the total triangle count.)

Normals

Per-vertex normals come from the surface where OCCT can give them (`Poly_Triangulation.Normal`, computed with `BRepLib_ToolTriangulatedShape` when the triangulation lacks them). This gives smooth shading on curved faces. If the surface cannot provide normals, BISCAD falls back to area-weighted face normals accumulated per vertex. Normals are flipped when the face orientation is `REVERSED`, so they always point out of the solid.

Edge polylines

B-rep edges are drawn as crisp lines, not as the boundary of triangles. BISCAD samples each edge into a polyline with `GCPnts_TangentialDeflection`, which places points so the chord never deviates from the true curve by more than the deflection, and the tangent never turns more than a set angle (0.2 rad here). A straight edge gets 2 points; a tight arc gets many. This is what gives the viewer and renderer their clean CAD look instead of faceted silhouettes.



The scene format, v1

The server sends the viewer one JSON document. The spec (from `tessellate.py`):

```
{
  "format": "biscad-scene", "version": 1, "units": "mm",
  "bbox": {"min": [x,y,z], "max": [x,y,z]},
  "parts": [{
    "id": "p0", "name": "bracket", "color": "#d8d6d0",
    "positions": <b64 float32 xyz>,    # vertex coordinates
    "normals": <b64 float32 xyz>,      # per-vertex normals
    "indices": <b64 uint32>,           # triangle vertex indices
    "faces": [{ "id": "p0/f3", "type": "plane", "start": 164, "count": 216,
      "area": 2743.27, "center": [...], "normal": [...] }],
    "edges": [{ "id": "p0/e7", "type": "line", "length": 37.0,
      "points": <b64 float32 xyz polyline> } ]
  }]
}
```

Field notes:

- **Geometry is base64 of raw typed arrays.** `positions` and `normals` are float32 triples; `indices` is uint32. The viewer decodes base64 straight into `Float32Array` / `Uint32Array` and hands them to three.js with no parsing (`viewer.js`, `b64ToArrayBuffer`, `f32`, `u32`). Base64 costs about 33% size over raw binary but rides inside one JSON response and gzips well.
- `faces` carries both the id and the triangle range, so picking and labelling need no extra lookup table.
- `type` is the OCCT geometry kind: `plane`, `cylinder`, `cone`, `sphere`, `torus`, `line`, `circle`, `bspline`. Curved faces also carry `radius` and `axis`; circles carry `center`.
- `topology.json` is a sibling document with the same descriptors but no mesh buffers. That is what the REST `/topology` endpoint and the MCP `get_topology` tool return, so an agent reads face geometry without downloading megabytes of triangles.

Size math

For one part, the raw buffer bytes are approximately:

```
positions:  V vertices × 3 × 4 bytes
normals:    V vertices × 3 × 4 bytes
indices:    T triangles × 3 × 4 bytes
edges:      S segments × 2 × 3 × 4 bytes    (both endpoints)
```

Base64 multiplies that by ~ 1.33 , and the face/edge descriptor JSON adds a few hundred bytes per face. Measured for the bracket at normal quality: 1,392 vertices, 460 triangles, 448 edge segments give a raw prediction of ~ 50 KB; the actual `scene.json` is 83 KB, the difference being base64 overhead plus the descriptor JSON (the bracket's 18 face descriptors are 2.6 KB, its 48 edge descriptors 5.9 KB). The scene gzips to about 18% of its size on the wire (the server runs GZip middleware).

Scenes range from tens of KB (a bracket) to a few MB (the 23-part quadruped at ~ 6 MB normal). Over gzip and an `immutable` cache header, this is fine for the web and cheap for an agent.

6. Headless rendering

An agent has no screen. A server has no GPU. But a multimodal model must *see* the part to check it, and the design report must embed pictures. So BISCAD ships a software renderer: `server/render.py`, Pillow and NumPy only, no OpenGL.

Why no GPU

GPU rendering on a server means a headless OpenGL context, a GPU in the box, and driver pain on Cloud Run and Hugging Face Spaces, which have neither. A pure-CPU renderer runs anywhere Python runs, in the same container, with no extra dependency. The cost is speed, and for a still image that cost is small: a normal part renders in 130 to 360 ms (Section 14). That is fast enough for an agent that is already waiting on the kernel.


The algorithm: painter's plus a depth image

The renderer is a classic **painter's algorithm** with one addition for correct edges.


1. **Project.** Build an orthographic camera from the requested view direction (`iso`, `front`, `top`, or an arbitrary `x,y,z`). Project every triangle vertex to 2D and keep its depth along the view axis. Scale so the model fills the frame with a small margin. Render at 2x supersampling, then downscale with Lanczos for clean edges.
2. **Shade.** Per triangle, compute a Lambert term from two lights plus a small Phong specular highlight, modulated by the part colour. Faces that are highlighted get painted international orange. Both sides are lit, so back faces seen through a hole are not black.
3. **Paint back to front.** Sort triangles by depth and draw them in order with Pillow polygons. The nearer triangle paints over the farther one. That is the painter's algorithm.
4. **Capture a depth image.** While painting the triangles, also paint their depth into a separate float image. This is the trick for edges.
5. **Draw visible edges only.** For each B-rep edge polyline, walk it in small steps. For each step, compare the edge's depth at that pixel against the depth image. Draw the segment only where the edge is at or in front of the surface (within a tolerance). This hides edges that are behind the solid, so you get clean hidden-line behaviour without a true hidden-line engine.
6. **Labels.** If requested, place each face's id at its projected centre, but only if that centre is not occluded (same depth test) and the face is large enough to label. A white box behind the text keeps it readable.
7. **A view triad** (X red, Y green, Z blue) goes in the corner so orientation is never ambiguous.

`render_grid` renders four views (iso, front, top, right) into one image. The MCP `render_grid` tool returns it: one picture, full understanding of a part.


ISO




FRONT



TOP



RIGHT



Limits

- **Per-triangle painting is $O(\text{triangles})$.** A 170k-triangle fine quadruped is slow to render. For agent use, normal quality (tens of thousands of triangles) is the right operating point.
- **No shadows, no ambient occlusion, no reflections.** The software renderer is for *reading* geometry, not for beauty shots. The web viewer (Section 12) does the pretty rendering on the client GPU.
- **Approximate edge visibility.** The depth test has a tolerance, so a few edges near a silhouette may flicker in or out. It is good enough to read the shape.

How agents use it

The MCP tools `render_view` and `render_grid` return real image blocks. A multimodal model gets the picture inside its tool result and checks: are there four legs? Does the hole go through? Is anything floating? The text-to-CAD loop (Section 11) feeds the agent a 4-view grid after every successful build and asks "does this match the request?". The renderer is the agent's eyes.

7. Sandbox security

BISCAD runs **untrusted Python**. Anyone with a key, and anonymous callers by IP, can POST a program that the server executes. The whole security model is defence in depth in `server/kernel.py`, and the known escapes are frozen as tests in `tests/test_api.py`.

Threat model

The attacker controls the full text of a Python program the server will `exec`. Their goals: read files (keys, other users' data), write files, run shell commands, open network connections, or crash/hang the server. BISCAD must allow `build123d`, `NumPy` and `math` to run freely while denying all of that.

The honest framing from `README.md`: this is best-effort in-process sandboxing. For a large public deployment, run the container under `gVisor` or `Firecracker` as well. The layers below raise the bar high; they are not a kernel-level jail.

Layer 1: AST guard

Before anything runs, `parse_and_check_script` walks the AST and rejects:

- **Imports not on the allowlist.** Only `build123d`, `math`, `cmath`, `random`, `itertools`, `functools`, `operator`, `typing`, `dataclasses`, `enum`, `copy`, `numpy`, `collections`, `statistics`, `fractions`, `decimal`, `string`, `re`. No `os`, `sys`, `subprocess`, `socket`, `pathlib`, `importlib`, `ctypes`. Relative imports are rejected outright.
- **Dunder attribute access.** Any `__class__`, `__globals__`, `__subclasses__` is a classic sandbox-escape primitive (walk from an object to `object.__subclasses__()` to arbitrary classes). All `__...` attribute access is banned.
- **Banned attribute names.** A denylist of I/O and introspection names: `export`, `import`, `write`, `read`, `save`, `load`, `open`, `system`, `popen`, `fork`, `globals`, `locals`, `vars`, `frame`, `mro`, `subclasses`, plus generator/traceback frame attributes (`gi_`, `cr_`, `tb_`, `f_`, `co_`) by prefix.
- **Banned builtins by name** as `Name` nodes: `open`, `exec`, `eval`, `compile`, `__import__`, `getattr`, `setattr`, `input`, `breakpoint`, `exit`.
- **Strings containing `__`.** This blocks `"{0.__class__}".format(x)` and `getattr`-by-string tricks, because the dangerous name never appears as a literal.

Layer 2: sanitized module copies

A script never sees a real module. `_sanitized_copy_of_module` builds a fresh module object holding only public, non-module, non-I/O members of the real module. So `build123d.os`, `typing.sys`, `re.enum.sys` do not exist in the sandbox: the chain to the real `os` is cut at the first hop. Sub-modules are dropped entirely (the copy holds no module-typed attributes), and I/O names like `export_step`, `Meshes`, `ExportSVG`, `import_brep` are stripped so a script cannot ask `build123d` to write a file. The guarded `__import__` returns these copies and nothing else.

This closed a whole class of escapes found in review (now tests): `build123d.exporters.os.getcwd()`, `import typing; typing.sys, numpy.save(...)`, `numpy.loadtxt('/etc/hostname')`.

Layer 3: no file writes during user code

Even with exports stripped, a library might try to write. While user code runs, `_execute_with_file_writes_blocked` sets `RLIMIT_FSIZE` to 0 and ignores `SIGXFSZ`, so any write of a non-empty file fails. The limit is restored to its real value immediately after, so BISCAD's own outputs (`scene.json`, `.brep`) write normally. This is the belt to the sanitized-module suspenders: writing is denied twice.

Layer 4: resource limits in the forked child

Every build runs in a `fork()`ed child (`run_in_forked_child`) that:

- Calls `setsid()` so the child is its own process group. On timeout the parent kills the whole group with `SIGKILL`, so a child that spawned helpers cannot leave orphans.
- Sets `RLIMIT_AS` (address space, default 3 GB), `RLIMIT_CPU` (seconds), and `RLIMIT_FSIZE` (512 MB ceiling for legitimate outputs).
- Is `daemon=True` and watched with a wall-clock timeout (30 to 180 s by plan). A hang is killed.

Fork isolation also means a segfault in OCCT (invalid geometry can crash the C++ layer) takes down only the child. The parent sees EOF on the pipe and returns "kernel process crashed", not a dead server.

Layer 5: read-only code in Docker

In the container (`Dockerfile`), the server code is copied as root and the process runs as an unprivileged user (`uid 1000`). Only `/data` is writable by that user. So even if a script reached the filesystem, it could not modify the server's own code.

The escapes found during review

These are real escapes caught before launch, each now a locked test (`test_sandbox_and_errors` in `tests/test_api.py`):

```
("import build123d\nbuild123d.exporters.os.getcwd()", "not available"),
("from build123d import *\nexport_stl(Box(1,1,1), '/tmp/x.stl')", "not available"),
("from build123d import *\nf = Box(1,1,1).export_step", "not available"),
("import numpy as np\nnp.save('/tmp/x.npy', np.zeros(3))", "not available"),
("import numpy as np\nnp.loadtxt('/etc/hostname')", "not available"),
("import typing\ntyping.sys", "not available"),
("g = (x for x in [1])\ng.gi_frame", "not available"),
("{0.__class__}'.format(1)", "'__'"),
("from build123d import *\nos", "not available"),
("import build123d.exporters", "not allowed"),
```

They cover: module-hopping to `os`, calling a stripped export, binding a method that would write, NumPy file I/O, sub-module traversal, generator-frame access, and format-string class walking. The test suite exists so no refactor silently reopens one.

Residual risks

- **In-process, same kernel.** A novel OCCT or CPython bug could still breach the interpreter boundary. The guard denies the known primitives, not unknown ones.
- **CPU/memory exhaustion up to the limits.** A script can burn its full CPU budget and memory ceiling. Rate limits and quotas (Section 13) bound the damage, and the worker semaphore bounds concurrency.
- **Side channels.** Timing and memory pressure are not addressed.

Recommendation for public scale: run the container under **gVisor** (a user-space kernel that intercepts syscalls) or **Firecracker** (a microVM per job). Either turns a breach of the Python sandbox into a breach of a throwaway sandboxed kernel, not the host. The in-process guard then becomes the first of three walls, not the only one.

8. Build-step capture

To replay how a part was built, BISCAD records each operation as it happens. The code is `server/steps.py`.

Wrapping the one funnel

Every build123d builder-mode operation that changes the part passes through one method: `Builder._add_to_context`. `StepRecorder.install()` monkey-patches that method with a wrapper. The wrapper:

1. Remembers the part object before the call.
2. Runs the real operation.
3. If the builder is a `BuildPart`, the mode is not `PRIVATE`, and the part object actually changed, records a step.

Because it hooks the one chokepoint, it catches primitives (`Box`, `Cylinder`), feature ops (`extrude`, `revolve`, `fillet`, `Hole`), and boolean modes alike. The recorder is installed only during the build and uninstalled after, so it never leaks into another job.

Frame walking: finding the op and the line

Knowing *that* the part changed is not enough; the step must say *what* ran and *where*. The recorder walks the Python call stack from the operation back toward the user's script:

- It climbs frames until it hits a frame whose filename is `<script>` (the user's code) and records that line number. That is the line shown in the report and the viewer.
- Along the way it looks for the frame whose `self` is a build123d `Shape` subclass (a `Box`, a `Hole`), which names the operation, or failing that the first non-private function name in the chain. From that frame's locals it pulls argument values (`amount`, `radius`, `length`, `angle`, `depth`, a fillet/chamfer edge `count`).

This is why a step can say "Drill a hole of diameter 6.5 mm" and point at line 17: the diameter came from the operation's locals, the line from the script frame.

Controlled-English descriptions


`describe()` turns the operation and its args into one or two short sentences, ASD-STE100 style: one action, active voice, units on numbers, a verb map (`extrude` → "Extrude", `offset` → "Shell", `Hole` → "Drill a hole"). It adds the volume change and the new face count. Measured output for the flange:

1. Add a cylinder of diameter 90 mm, height 8 mm. Volume increases by 50,894 mm³. The part has 3 faces.
2. Add a cylinder of diameter 46 mm, height 28 mm. Volume increases by 33,238 mm³. The part has 5 faces.
3. Fillet the edges with radius 4 mm. Volume increases by 515 mm³. The part has 6 faces.
4. Drill a hole of diameter 22 mm, through the part. Volume decreases by 10,644 mm³. The part has 7 faces.
5. Drill a counterbored hole of diameter 6.6 mm, through the part. Volume decreases by 2,737 mm³. The part has 25 faces.
6. Chamfer 2 edges with size 1 mm. Volume decreases by 107 mm³. The part has 27 faces.


A person reads that like a recipe and knows exactly what the agent did.

Draft meshing of snapshots


`finalize()` keeps one entry per step that actually changed volume or face count (it drops no-op or duplicate records). For each kept step, the kernel writes a **light** directory: a scene and a draft mesh, no `.brep` (`_write_build_step_snapshots`, quality forced to `draft`). Draft because there can be many steps (the quadruped has 60), and each is a preview the viewer's steps timeline or the explainer GIF will show small.




1. extrude (line 11)




2. fillet (line 13)




3. hole (line 17)



4. hole (line 22)



5. chamfer (line 23)



6. fillet (line 24)

Step meshing is not free: it was ~20% of the quadruped's build time (Section 14), because meshing 60 intermediate solids is 60 extra tessellations. It is capped at 80 steps (`MAX_STEPS`) to bound the cost.

9. Analysis algorithms

Analysis runs on a saved version: it loads the `.brep` files inside a forked child and computes. All of it is in `server/analysis.py`. For each operation below: the math, the cost, and how it fails.

measure — distance, angle, parallelism

Math. For two entities (`p0/f3`, `p0/e7`, a vertex, or a whole part), the minimum distance uses OCCT's `BRepExtrema_DistShapeShape`, which finds the closest points between two shapes exactly (not on the mesh, on the B-rep). Angle comes from the dot product of the two entities' directions: a plane's normal, a cylinder's axis, a line's tangent, a circle's normal. `parallel` and `perpendicular` fall out of the cosine.

Cost. `BRepExtrema_DistShapeShape` is roughly $O(F_a \times F_b)$ in the sub-faces of the two entities, but entities are small (one face, one edge), so it is milliseconds. Measured: ~60 ms including the fork and `.brep` load.

Failure modes. An unknown ref string raises a clear error. Entities with no direction (a sphere face, a spline) give a distance but no angle.


The viewer computes a fast approximate measure on the client for instant feedback, then calls this endpoint and replaces it with the exact kernel answer (Section 12).

section — a half-space boolean

Math. To cut the model with a plane, BISCAD intersects every part with a very large half-space box (1e5 mm) positioned on the plane: `kept = shape & half`. The cut faces are the faces of the result that lie in the plane (centre on the plane, normal parallel to the plane normal). Their total area is the section area. An SVG is drawn by projecting those faces to the plane's local 2D frame and filling them with a hatch pattern.

Cost. One boolean per part, so $O(\text{parts})$ booleans; a boolean is the expensive operation. Measured ~33 ms for the one-part flange.

Failure modes. A boolean that fails on a bad solid skips that part (caught). A plane that misses the model gives zero area and an empty SVG.



Section $y = 0$: 4 regions, cut area 896.9 mm² (highlighted).

mass — GProp and the inertia tensor

Math. OCCT's `BRepGProp.VolumeProperties_s` integrates over the solid to give volume, centre of mass, and the inertia matrix about the centre of mass. BISCAD multiplies by the density (g/cm³, converting mm³ → cm³) to get mass in grams, and scales the inertia into **kg·mm²**. For an assembly it sums masses and computes the combined centre of mass as the mass-weighted mean.

Cost. One volume integration per part. Measured ~5 ms for the flange. Verified against the summary: mass equals volume/1000 × density to within rounding.


Failure modes. An open or invalid shell gives a meaningless volume; the validity flag warns you. Units are the usual trap, so the response labels every field (`mass_g`, `inertia_kg_mm2_about_com`).

DFM checks — overhang, min wall, tiny edges

`check(process)` runs design-for-manufacturing rules for `fdm`, `cnc` or `sheet`.

- **Build envelope.** Compare the sorted bounding-box dimensions to the process envelope (FDM 256³, CNC 600×400×300, sheet 1500×3000×50 mm). Error if it does not fit.
- **Validity and disconnected solids.** Report invalid solids and parts made of multiple disconnected lumps.
- **Tiny edges.** Any edge shorter than 0.2 mm is flagged (often a sliver from a failed boolean; hard to manufacture). O(edges).
- **Overhangs (FDM).** For each face, if its normal points more than 45° below horizontal (`normal.Z < -0.707`) and it is not on the build plate, it needs support. BISCAD reports the face refs and the overhang area fraction. O(faces).
- **Min wall estimate (ray casting).** From each face centre, shoot a ray inward along `-normal` with `BRepIntCurveSurface_Inter` and take the nearest exit distance. The smallest over all faces is the estimated minimum wall thickness. Capped at 400 faces for cost. This is an estimate, not a medial-axis thickness, but it catches thin walls cheaply. O(faces × intersection cost).

Measured: on the enclosure, the check finds the correct overhanging vent faces and estimates a 2.4 mm minimum wall. On the whole quadruped (23 parts) the check takes ~3.2 s because it is 23 parts × hundreds of faces of ray casting.



FDM check: 8 faces overhang more than 45 deg (highlighted). View from below.

Failure modes. Ray casting from a point exactly on a surface can miss; a 1e-3 mm back-off handles it. The overhang test assumes the part prints in its modelled orientation.

diff — geometric comparison of two versions

Math. Given versions A and B, BISCAD computes $B - A$ and $A - B$ as booleans to get volume added and removed, and compares face *signatures* (type, rounded area, rounded centre) as sets to count faces added and removed. `identical` is true when both signature sets match and volumes agree.

Cost. Two booleans plus set operations. Measured ~60 ms for the bracket. The signature comparison is a heuristic: it can be fooled by two different faces with the same area and centre, but in practice it is a good, cheap diff.

Failure modes. A boolean that fails records `boolean_error` and still returns the face/volume comparison.

interference — clash detection

Math. For every pair of parts in an assembly, first test their bounding boxes. If the boxes do not overlap, the parts cannot clash, so skip (the **bbox prefilter**, $O(n^2)$ cheap tests). Only for boxes that overlap does BISCAD compute the actual common volume $a \& b$ (the expensive boolean) and report pairs whose overlap exceeds a tolerance, sorted by volume, with the clash centre.

Cost. $O(n^2)$ bbox tests plus one boolean per candidate pair. The prefilter is what makes it affordable: measured on the 23-part quadruped, only 18 of 253 pairs survive the prefilter, and the whole check is ~2.25 s. Without the prefilter it would be 253 booleans.

Failure modes. Touching-but-not-overlapping parts may register a tiny overlap from tolerance; the tolerance threshold filters those.

BOM — grouping identical parts

Math. Group parts by a signature of (rounded volume, rounded area, face count, sorted bounding-box size). Parts with the same signature are treated as the same item and counted. Items are sorted by quantity.

Cost. $O(\text{parts})$, plus the per-part property reads. Measured ~ 1 s for the quadruped, dominated by reading volume/area/faces for 23 parts. Verified: it groups the quadruped's 23 parts into 8 unique items, and in a test with three identical boxes it reports quantity 3.

Failure modes. The signature is a heuristic. Two genuinely different parts with the same volume, area, face count and bounding box would be merged. Mirror-image parts (left/right) have the same signature and are grouped, which is usually what a BOM wants but not always.

10. Exports, imports and drawings

3D exchange formats

`analysis.export(vdir, fmt)` loads the `.brep`, assembles a compound (named, coloured), and writes the requested format, caching the file so a second request is instant. Supported: `step`, `stl`, `glb`, `3mf`, `brep`, `obj`, plus the 2D `svg/dxf` drawings.

- **STEP** (`export_step`) is the exact B-rep interchange format: the right choice to move the model to another CAD system. Measured ~13 ms, 75 KB for the flange.
- **STL / 3MF / GLB** are meshes for printing and web: tessellated at a fixed tolerance. 3MF carries colour and units; GLB is the web/AR mesh. 3MF is slower (~226 ms) because it re-meshes through OCCT's `Mesh`.
- **BREP** is OCCT's native dump, fastest of all (~1 ms) and the format BISCAD stores internally.
- **OBJ** is written directly from the cached `mesh.npz`, so it needs no kernel work.

Imports

`POST /v1/import` accepts STEP, STP, BREP and STL. The file is saved, then a forked child imports it (`import_step` / `import_brep` / `import_stl`), tessellates it into the same scene format, and stores it as a new document's first version. A round-trip is tested: export the flange to STEP, re-import it, and the face count matches.

Drawings: HLR projection to SVG and DXF

A 2D engineering drawing needs the part's silhouette and visible edges projected flat, with hidden edges dashed. That is **hidden-line removal (HLR)**, and OCCT does it.

`_drawing()` lays out four views (front, top, right, iso) on a sheet. For each, it calls `shape.project_to_viewport(...)`, which runs OCCT's HLR algorithm and returns two edge sets: visible and hidden. BISCAD writes the visible edges to a "visible" layer and the hidden edges to a dashed "hidden" layer, via `build123d's` `ExportSVG` / `ExportDXF`. SVG is for viewing, DXF for CAM and laser cutters.

HLR is the most expensive export because it is a full projection-and-occlusion pass per view. Measured on the flange: SVG ~690 ms, DXF ~400 ms. That is the price of a real drawing with correct hidden lines, and it is still well under a second.

11. Text-to-CAD agent loop and MCP tool design

The loop: write, build, look, fix

`server/agent.py` is the text-to-CAD feature. It needs `ANTHROPIC_API_KEY` on the server; without it, `/v1/agent` returns 501 and the studio's Ask box is disabled. The loop is BISCAD's own, not a generic tool runner:

```
sequenceDiagram
    participant U as User prompt
    participant M as Claude (model)
    participant K as Kernel
    U->>M: "a 20 mm plate with a hole"
    loop up to MAX_ROUNDS (4)
        M->>M: write a build123d program
        M->>K: build it (a real version)
        alt build failed
            K-->>M: error text + logs → "Fix it."
        else build ok
            K-->>M: 4-view render grid + facts (bbox, volume, steps)
            M->>M: does this match the request?
            M-->>U: DONE (or a corrected program → loop)
        end
    end
end
```

Design decisions:

- **Every round is a real version.** The agent's attempts are saved versions with parent links, so the user can scrub through the agent's thinking, not just see the final answer. The studio shows round chips (R1, R2, ...) you can click.
- **Render-and-fix.** After a good build, the agent is shown a 4-view grid (the software renderer, Section 6) and the build steps and asked to verify against the request. This is the Karpathy loop made literal: the model checks its own output visually.
- **Errors feed back.** A failed build sends the error and the last 1,500 characters of logs back with "Fix it." A build123d error names the script line, so the model edits the right place.

Cost control and quota gating

Each agent run spends the server owner's Anthropic credits, so it is gated hard:

- **Auth required.** Anonymous callers get 401. Text-to-CAD needs a key.
- **Per-plan monthly cap.** `agent_runs_month` (free plan default 10, from `BISCAD_FREE_AGENT_RUNS`). Tracked in a separate usage bucket (`agent:<owner>`), so agent runs do not eat the build quota and vice versa. Over cap → 402.
- **Bounded rounds.** `MAX_ROUNDS` (default 4) caps the number of model turns, so one prompt cannot loop forever and drain credits.

Fallbacks

The model call (`_ask`) asks for medium effort and passes server-side fallbacks, so a transient overload of the primary model falls back rather than failing. API errors map to clean HTTP codes: rate limit → 429, bad credentials → 501, refusal → 422, connection failure → 502. The studio shows a human message for each.

MCP tool design for agents

The MCP server (`server/mcp_server.py`) mirrors the REST API as tools, but the shape of the tools is tuned for a model, not a human:

- **Images in results.** `build_model`, `render_view` and `render_grid` return real image blocks, so a multimodal agent sees the geometry inside the tool result. A build comes back as a compact JSON summary *plus* a labelled iso render.
- **Compact JSON.** `_brief()` strips the summary to what matters (bbox, volume, mass, per-part validity, build steps as one-liners, a viewer URL, a report URL) and serialises with no whitespace. Small results keep the agent's context cheap.
- **One guide text, shared.** `GUIDE` is a single document: the allowed imports, the `result = / show(...)` contract, the `params = {...}` rule, the `p0/f3` reference scheme, and a build123d cheat sheet. It is the MCP server's `instructions`, and the text-to-CAD system prompt prepends it, so the agent learns the rules once and the REST and MCP paths teach the same thing.
- **Stable ids everywhere.** Every tool speaks `p0/f3`. `get_topology` returns ids and geometry; `measure`, `section`, `check_manufacturability` take ids; `render_view` highlights ids. The agent points, the human sees the same highlight.

The workflow the guide recommends: `build_model` → look at the render → `get_topology` to find ids → `measure` / `section` / `check` → `export_model`.

12. Web viewer and studio internals

The viewer is `web/js/viewer.js`, a reusable three.js ES module (`new Viewer(container, opts)`). It is the biggest single file in the project (~2,270 lines) and the one that makes the geometry feel like a real CAD viewer. three.js and CodeMirror are vendored under `web/vendor/` and loaded via an import map; there is no build step.

Scene construction from base64 buffers

`loadScene` decodes the scene (Section 5). For each part, `_buildPart`:

- Decodes `positions`, `normals`, `indices` straight from base64 into typed arrays and sets them as `BufferAttributes` on a `BufferGeometry` with zero copying or parsing.
- Builds a `MeshPhysicalMaterial` tuned per part colour (dark parts get a little clearcoat and metalness so the quadruped's black motors read well).
- Decodes edge polylines into a `LineSegments2` (fat lines) so B-rep edges stay crisp at any zoom.
- Up axis is Z (`THREE.Object3D.DEFAULT_UP.set(0,0,1)`), matching CAD convention.

GPU id picking

Clicking a face must return `p0/f3` exactly. The viewer uses **GPU colour picking**:

1. Each triangle gets a per-vertex `pickId` attribute: an integer naming its face (assigned from the same face ranges as the scene; `_buildPart` fills `pick[]` over each face's `start..start+count`). Edges get an instanced `instanceEdgeId`.
2. A tiny custom shader (`PICK_VS` / `PICK_FS`) encodes that integer id into an RGB colour (`vcEncode`: id split across R, G, B channels).
3. On click, `_pickAt(x, y)` swaps every part to the pick material, sets a 1x1 camera view offset at the cursor, renders into a 1x1 render target, and reads back the one pixel with `readRenderTargetPixels`.
4. The pixel's RGB decodes to the id, which indexes a lookup table (`_pickTable`) back to `{part, kind, index}` and then to `p0/f3`.

This is exact (no raycast math), works for faces and fat-line edges, and is one pixel of readback per click. A `filter` lets the studio restrict picking to faces, edges or whole parts.

Section: clipping plane plus back-face caps

Sectioning uses three.js **local clipping planes**. `_applySection` builds one `THREE.Plane` from the section normal and offset and pushes it into the shared `clipPlanes` array that every material references. The GPU clips away everything on one side.

A raw clip leaves the solid looking hollow. So each part also has a **cap mesh**: the same geometry drawn with `side: THREE.BackSide` and a flat hatched shader (`onBeforeCompile` injects a diagonal-line pattern). The back faces, seen through the cut, paint the interior as a solid hatched cross-section, so a sectioned part looks filled, like an engineering section view. The cap is hidden in x-ray mode and during picking.

Changing the number of clip planes forces a material recompile, which is why `_applySection` flips `needsUpdate` on every affected material.

Measure: fast client, exact kernel

When you pick two entities, `_runMeasure` computes a measurement immediately on the client: centre-to-centre, parallel-plane distance, axis-to-axis, line-to-line closest approach (`segSegClosest`, a standard segment-segment closest-point solver), face-to-edge distance and angle. This is instant feedback. Then, if a `measureProvider` is wired (the studio passes `/v1/versions/{id}/measure`), it calls the kernel for the **exact** `BRepExtrema` distance and replaces the client estimate, tagging the result `source: "kernel"`. Fast first, exact second.

Explode

`_computeExplodeDirs` gives each part an explosion direction: the vector from the model centre to the part centre. Parts near the centre get a fanned-out radial direction so they do not pile up. `_applyExplode` scales each part along its direction by the explode amount; `animateExplode` eases it over 600 ms. The toolbar slider and the `X` key drive it. Explode is hidden for single-part models.

Other viewer features

- **Render modes:** shaded+edges, shaded, hidden-line, wireframe, x-ray.

- **Contact shadows:** two depth-rendered, blurred shadow layers (tight contact plus broad ambient) for a grounded look, all on the client GPU.
- **Optional GTAO** ambient occlusion, lazy-loaded from three addons, with fat lines and helpers excluded from the AO buffer.
- **View cube, infinite fading grid** (a shader that fades with distance and draws coloured X/Y axes), **bounding box, screenshot.**
- **StepsBar:** replays build steps, ghosting the previous step behind the current one so added/removed material reads clearly. It loads each step's light scene on demand from its `scene_url`.

Embed mode

`view.html?v=<id>&embed=1` strips the chrome (toolbar, panels, grid, steps), makes the background transparent, disables picking unless `&pick=1`, and can auto-rotate. The studio's share dialog produces the `<iframe>` snippet. `?api=` overrides the backend origin, so a viewer hosted anywhere can point at any BISCAD server.

The studio

`web/js/studio.js` wires a CodeMirror editor, a parameter panel, the viewer and version history against the v1 API.

- **Live parameters.** It reads `params = {...}` from the script locally for instant sliders, then the server's `param_schema` wins. Moving a slider debounces and rebuilds (`scheduleBuild`), so the model updates as you drag.
- **Build on demand** with a race guard (`buildSeq`): a stale build result is dropped if a newer build started.
- **Offline demo.** If the API is unreachable, the studio loads sample scenes from `web/samples/*.json`, so the viewer works with no backend.
- **Ask box** for text-to-CAD, with round chips and a live "round N · Ns" status.
- **Agent view modal** shows exactly what an agent sees: the labelled render URL.

13. Data model, accounts, quotas, versions

Store

`server/store.py` is SQLite (WAL mode) plus a files tree under `data/versions/<vid>/`. Four tables:

- **keys** — API keys (SHA-256 hashed, never stored raw), email, plan.
- **documents** — owner, name, public flag, head version.
- **versions** — the build record: script, params, status, error, logs, summary, param_schema, steps, parent, owner. The geometry itself lives in the files tree, not the DB.
- **usage** — per owner per month: builds, calls, compute_ms.

Accounts and the capability model

- **API key.** `POST /v1/keys {email}` returns `bsc_...`. The key hashes to a row; the raw key is shown once. `BISCAD_ADMIN_KEY` seeds an unlimited key for the owner.
- **Anonymous = IP bucket.** No key means the `anon` plan, metered by IP, with a small playground quota. This lets the landing page and embeds work with no sign-up.
- **Version id = capability link.** A version id is unguessable (`secrets.token_urlsafe`). Anyone holding it can view, render and export that version: that is how share links and embeds work without auth. **Documents**, by contrast, are listed and modified only by their owner's key (or are explicitly `public`). The test suite checks that another key gets 404 trying to add a version to your document.

Quotas and rate limiting

`server/core.py` decides, for every request, before any work:

- **Rate limit** — a token bucket per owner, refilled at the plan's requests-per-minute. Over → 429.
- **Build quota** — builds per month and compute-ms per month, per plan. Over → 402 with a message pointing to a free key.
- **Concurrency** — a bounded semaphore (`BISCAD_WORKERS`, default = CPU count) caps how many kernel children run at once. Over the wait → 503.

Plans (from `store.py`):

Plan	builds/mo	compute s/mo	req/min	build timeout	agent runs/mo
anon	300	600	30	30 s	0
free	2,000	3,600	120	60 s	10
pro	100,000	200,000	600	180 s	1,000
unlimited	huge	huge	huge	600 s	huge

The Pro/Scale tiers on the landing page are copy only. No billing is wired; nothing was purchased. The plans exist in code so the quota machinery is real and testable.

Versions as capabilities and branches

A document is a tree of versions. Each version has a `parent`, so you can branch from any point (not just the tip). `new_version` defaults the parent to the document head but accepts any parent. A version reuses its parent's script if you send only new params, so "same design, width 80" is one small call. The geometric `diff` (Section 9) compares any two versions. This is configurations and branching without a feature-tree UI: it is just builds with parents.

14. Performance: measured numbers

All numbers below were measured on this machine for this handbook: a 4-vCPU Intel Xeon at 2.10 GHz, 15 GB RAM, build123d 0.13.0 on OpenCascade, Python 3.11. Builds were run through the in-process kernel (`server/kernel._build_job`) so the timings are the real mesh-and-capture cost, not network.

Build time, mesh size, render time per example

At **normal** quality, with the renderer timings for the iso view (800×600), a labelled iso, and the 4-view grid:

Example	Parts	Faces	Edges	Triangles	Steps	exec ms	total ms	scene KB	scene gz KB	render ms	render+labels ms	grid ms
bracket	1	18	48	460	6	102	308	81	15	143	133	278
flange	1	27	51	1,926	6	126	373	267	62	249	229	596
enclosure	1	83	200	2,249	5	271	759	345	68	241	233	533
gear	1	392	1,170	926	3	457	2,141	381	63	246	265	512
leg	5	71	129	7,565	14	266	870	869	249	356	357	1,239
quadruped	23	1,128	2,616	50,870	60	2,553	9,692	6,188	1,818	1,873	1,913	6,459

Notes:

- **exec** is the time running the user's build123d program (the booleans, fillets, holes). **total** is exec plus tessellation, step-snapshot meshing, and summary.
- A **simple part is sub-second end to end**. Bracket 0.31 s, flange 0.37 s. This is the common case for an agent, and it is the number the whole pricing argument rests on.
- The **gear** has many teeth, so it has 392 faces and 1,170 edges even though its triangle count is modest: exec is dominated by building the involute profile and the tooth geometry, not meshing.
- The **quadruped** (23 parts, 60 steps) is the stress case at ~9.7 s. It is a full robot assembly, not a typical agent iteration.

Triangle budget across quality

Example	draft tris	normal tris	fine tris
bracket	244	460	929
flange	834	1,926	5,128
enclosure	969	2,249	6,612
leg	2,583	7,565	23,264
quadruped	17,716	50,870	169,460

Fine quality roughly 2–3x's the triangles of normal. Draft roughly halves them. The quadruped at fine produces an 18 MB scene; normal (6 MB) is the right default.

Where the time goes: a profiled heavy build

`cProfile` on the quadruped build (23 parts, 60 steps, in-process, ~12 s under the profiler) ranks the cumulative cost like this:

Phase	Share of build	What it is
Final mesh (<code>write_shape_dir</code> → <code>build_scene</code> → <code>mesh_shape</code>)	~2.2 s	Tessellating the 23 final parts into the scene and buffers
Step snapshots	~1.9 s	Draft-meshing the 60 intermediate solids
Summary (<code>summarize_parts</code>)	~2.3 s	Per-part volume, area, bbox, validity
The user program (<code>exec</code>)	~2.6–2.9 s	build123d operations
JSON serialisation	~1.4 s	Writing the scene/topology/step JSON

By single C++ function (`totttime`):

- `BRepAlgoAPI.Build` (~1.47 s) — the boolean engine. Every union and cut.
- `is_valid` (~1.28 s) — B-rep validity checking, called per part in the summary.
- `_mesh_face` (~1.72 s cumulative, ~0.82 s own) — pulling triangles out of OCCT one node at a time in Python.
- `BRepBndLib.AddOptimal_s` (~0.86 s) — optimal bounding boxes.
- `BRepGProp.SurfaceProperties_s` / `VolumeProperties_s` (~0.7 s) — area and volume integration.

The lesson: for a heavy assembly, time is split roughly evenly between **computing the geometry** (booleans), **checking it** (validity, bounding boxes, mass properties), and **turning it into triangles and JSON** (meshing, serialisation). The per-vertex, per-triangle Python loop in `_mesh_face` is the one piece of BISCAD's own code that could be vectorised to claw back time; the OCCT calls are irreducible.

For a simple part none of this matters: the whole thing is a few hundred milliseconds and dominated by the one-time cost of the operations themselves.

Fork overhead

A bare forked round-trip (`run_in_forked_child(lambda: 1)`) is ~14 ms. A full box build through the fork path (guard, fork, build, mesh, write, return) is ~48 ms for 5 builds, so ~10 ms each of real work on top of fork. Because `build123d/OCCT` is preloaded in the parent, the fork is cheap; a fresh `spawn` would re-import OCCT and cost over 2 s (the measured import time of `build123d` alone was 2.1 s).

15. Deployment and scaling

One Dockerfile serves five deployment paths (`README.md`, `deploy/`).

Docker

The image (`Dockerfile`) is `python:3.11-slim` plus the OCCT runtime system libraries (`libgl1`, `libglu1`, `libgomp1`, `fontconfig`, `DejaVu` fonts for render labels). The server code is copied as root; the process runs as `uid 1000`; only `/data` is writable. `uvicorn` serves REST + MCP + static web on one port. A `HEALTHCHECK` hits `/v1/health`.

Cloud Run autoscaling, with caps

`deploy/cloudrun.sh` deploys to Google Cloud Run: 2 vCPU, 4 GiB, concurrency 8, **scales 0** → `BISCAD_MAX_INSTANCES` (**default 2**). Scaling to zero means an idle service costs nothing. The hard max-instance cap is the money guard rail required by `AGENTS.md` ("autoscaling deploys must cap max instances"). The script also sets a **\$5 budget alert** if it finds a billing account. The in-app rate limits and quotas bound per-caller load inside each instance.

Hugging Face Spaces

`deploy/huggingface.sh` pushes to a free Space (Docker SDK, 2 vCPU / 16 GB, no card). Permanent free hosting. The catch, stated plainly: Space storage is ephemeral, so documents reset on restart unless you point `BISCAD_DATA` at persistent storage.

Cloudflare quick tunnel

`deploy/public.sh` runs BISCAD locally (Docker or a `venv`) and opens a free Cloudflare quick tunnel, printing a public `trycloudflare.com` URL with no account. This is the "share it from my laptop in 30 seconds" path.

Cost model: per build vs per Onshape call

Cloud Run request-based pricing in a Tier 1 region is **\$0.000024 per vCPU-second** and **\$0.000025 per GiB-second**, with a monthly free tier of **180,000 vCPU-seconds** and **360,000 GiB-seconds** (Google Cloud Run pricing). Take a typical agent build: ~0.3 s of wall time on 2 vCPU and 4 GiB.

```
CPU:    0.3 s × 2 vCPU × $0.000024 = $0.0000144
Memory: 0.3 s × 4 GiB × $0.000025 = $0.000030
Total ≈ $0.000017 per build (about 1/60,000 of a dollar)
```

The free tier alone covers roughly $180,000 / (0.3 \times 2) \approx$ **300,000 builds per month** at no cost. Even paid, a build is well under a hundredth of a cent.

Compare Onshape. A Professional seat is \$2,500 per year for 5,000 API calls, so the *marginal* cost once you hit the wall is effectively "buy another \$2,500 seat", i.e. **\$0.50 per call** amortised over the allowance, and then a hard 402 stop. BISCAD's build is roughly **four orders of magnitude cheaper** than an Onshape API call, and it does not stop at an annual wall. That gap is the entire reason the project exists (Section 1).

16. Testing strategy

Tests are `tests/test_api.py`, run with `python -m pytest -q tests`. They are end-to-end against the real app via FastAPI's `TestClient`: real kernel, real store, real renderer, real MCP. 41 tests, all green (re-verified for this handbook against a snapshot of the server code).

What they cover:

- **Every example builds, is valid, and captures steps.** Parametrised over all six examples. This catches regressions in `build123d` usage, tessellation and step capture at once.
- **Scene format invariants.** First face is `p0/f0`; the base64 index buffer length equals the summed triangle count $\times 3 \times 4$ bytes. This locks the on-wire contract.
- **Parameter override and coercion**, including a string `"30"` coerced to int.
- **Sandbox escapes as locked tests** (Section 7). The known escapes are frozen so a refactor cannot silently reopen one. This is the most important test group for a service that runs untrusted code.
- **All exports produce non-trivial files**; STEP round-trips through import with the face count preserved.
- **Render, grid, explainer GIF, report** all return images / HTML.
- **Every analysis op** (measure, section, mass, check, topology filter, diff, interference, BOM) with value assertions (e.g. the three-box clash volume is 500 mm³; BOM groups three identical boxes as quantity 3).
- **Documents, versions, branching, ownership** (another key gets 404).
- **Auth and quotas** (bad key $\rightarrow$ 401, anonymous flagged).
- **MCP protocol**: initialize, tools/list contains the core tools, a `build_model` call returns JSON plus an image block.
- **The agent loop with a fake model**: a monkeypatched `_ask` returns a bad program, then a good one, then `DONE`, and the test asserts the error was fed back and the render image was shown. This tests the loop without spending credits.

The philosophy: test the contract at the edge (HTTP in, JSON/images out), run the real geometry kernel, and freeze every security escape as a permanent test. There are no mocks of the kernel; a passing suite means real parts really built.

17. Roadmap

BISCAD matches Onshape's API on build, reference, analysis, export and versioning, far more cheaply. Onshape still has four things BISCAD lacks. Here is each, and how to add it, mostly in few lines, because the model is code and the kernel is OCCT.

Sketches with a constraint solver

Onshape sketches are dimensionally constrained: you draw roughly, add "this edge is 10 mm, these two are parallel, this is tangent", and a solver places the geometry. `build123d` sketches are imperative (you give coordinates). To add constraints: wrap a 2D geometric constraint solver (OCCT has `PlaneGCS`-style building blocks, or vendor FreeCAD's `planegcs`) and expose a small sketch DSL that emits solved coordinates into a `build123d` `BuildSketch`. The build-step and reference machinery already handle whatever geometry comes out. Estimate: a new `sketch.py` plus one MCP tool.

Assembly mates

Onshape assemblies position parts by mates (coincident, revolute, slider) with degrees of freedom. BISCAD assemblies are parts placed by explicit `Location`. To add mates: define mate connectors as named frames on parts (a face id plus an axis, which the reference system already names), solve the transforms that satisfy the mate set, and apply them as `Locations` before building the `Compound`. Interference checking (Section 9) already validates the result. Estimate: a `mates.py` solver over existing references.

PDM and release management

Onshape has part numbers, revisions, and release workflows. BISCAD has documents and versions. To add release: a `released` flag and a revision label on a version, an immutable copy on release, and a BOM export (already built) tagged with revisions. This is mostly store and API work, no kernel change. Estimate: a few columns and endpoints.

Real-time collaboration

Onshape has multiple cursors in one document. BISCAD is single-writer per document. To add collaboration: a WebSocket channel per document broadcasting version events and viewer selections, and optimistic version creation. The viewer already emits selection and state events; the capability-link model already lets many clients view one version. Estimate: a WebSocket endpoint and a thin client channel.

The theme: because a BISCAD model is a program and the kernel is OCCT, most of what Onshape's UI provides is a small amount of code on top of machinery BISCAD already has.

18. Glossary

- **ASD-STE100** — Simplified Technical English, the controlled-English style used for BISCAD prose and step descriptions: short active sentences, one idea each, units on numbers.
- **B-rep** — boundary representation. A solid stored as exact trimmed surfaces and curves stitched into a watertight shell. See Section 3.
- **Boolean** — union, cut or intersect of two solids. The core and costliest B-rep operation.
- **BRepMesh** — OCCT's incremental surface mesher. Turns B-rep faces into triangles.
- **build123d** — the Python CAD library BISCAD runs; a code-first wrapper over OCCT.
- **Capability link** — an unguessable id that grants access by possession. BISCAD version ids work this way.
- **Chamfer / Fillet** — replace an edge with a flat bevel / a rounded blend surface.
- **CSG** — constructive solid geometry; a boolean tree of primitives.
- **Deflection** — the allowed gap (linear, mm) or angle (angular, rad) between a mesh and the true surface. Controls mesh density.
- **DFM** — design for manufacturing. BISCAD's `check` finds overhangs, thin walls, tiny edges, over-size parts.
- **GProp** — OCCT's global properties: volume, area, centre of mass, inertia.
- **HLR** — hidden-line removal. Projecting a 3D solid to 2D with hidden edges dashed, for drawings.
- **MCP** — Model Context Protocol. The tool interface agents call BISCAD through.
- **OCCT / OpenCascade** — the open-source B-rep geometry kernel under build123d.
- **Painter's algorithm** — draw triangles back to front so nearer ones cover farther ones. BISCAD's software renderer.
- **Scene (biscad-scene v1)** — the JSON document the server sends the viewer: mesh buffers plus face/edge descriptors with stable ids. Section 5.
- **TopoDS** — OCCT's topology data structures (Compound → Solid → Face → Edge → Vertex).
- **Topological naming problem (TNP)** — face/edge ids shift when a model's structure changes, because ids come from regeneration order. Section 4.
- **Tessellation** — converting exact B-rep into triangles.
- **Version** — one built result (script + params → solids + files). The unit of storage, sharing and diff.

19. Reading list

Geometry kernel and B-rep

- OpenCascade Technology documentation — Modeling Data and Modeling Algorithms: <https://dev.opencascade.org/doc/overview/html/>
- OCCT `BRepMesh` (incremental meshing) reference and user guide (meshing chapter of the OCCT docs).
- OCCT `BRepExtrema`, `BRepGProp`, `BRepAlgoAPI` class references (the exact classes BISCAD uses for measure, mass and booleans).
- Mäntylä, M. *An Introduction to Solid Modeling*. Computer Science Press, 1988. The classic text on B-rep and Euler operators.
- Hoffmann, C. *Geometric and Solid Modeling*. Morgan Kaufmann, 1989.

build123d

- build123d documentation: <https://build123d.readthedocs.io> — builder vs algebra mode, selectors, the operation reference.
- CadQuery (build123d's predecessor in spirit): <https://cadquery.readthedocs.io>

Topological naming

- Onshape's approach to deterministic/persistent ids — Onshape blog and documentation on the topological naming problem and why feature references survive edits.
- FreeCAD "Toponaming" / topological naming mitigation (merged for FreeCAD 1.0), and Realthunder's element-mapping work: https://wiki.freecad.org/Topological_naming_problem
- Kripac, J. "A mechanism for persistently naming topological entities in history-based parametric solid models." *Computer-Aided Design*, 1997 — the foundational paper on persistent naming.

Rendering and meshing

- Foley, van Dam, et al. *Computer Graphics: Principles and Practice* — the painter's algorithm, projection, shading models.
- Morten S. Mikkelsen. "Simulation of Wrinkled Surfaces Revisited" (2008) and his notes on consistent tangent-space / hex-tiling — on generating stable, aliasing-free surface detail and tiling without repetition, relevant to texture and detail work beyond flat shading.
- three.js documentation and examples: <https://threejs.org/docs/> — clipping planes, render targets, `LineSegments2`.

Agents and oversight

- Andrej Karpathy on working with LLMs and keeping humans in the loop (talks and writing on AI as tools people must be able to verify) — the motivation for renders, steps and reports.
- Model Context Protocol specification: <https://modelcontextprotocol.io>

Manufacturing and formats

- ISO 10303 (STEP) overview — the B-rep interchange standard.
- 3MF Consortium specification: <https://3mf.io> — colour, units, mesh for printing.

Service limits and cost

- Onshape API limits: <https://onshape-public.github.io/docs/auth/limits/>
- Onshape pricing: <https://www.onshape.com/en/pricing>
- Google Cloud Run pricing: <https://cloud.google.com/run/pricing>

End of handbook. The source of truth is the code in `server/`, `web/js/` and `client/`; where this document and the code disagree, the code wins. Figures in `docs/img/` were generated by `server/render.py` from in-kernel builds of the example programs.